

HiLo-STeP++: LOW-RANK STOCHASTIC TAYLOR PROBING FOR HIGH-DIMENSIONAL DIFFUSIONS

Anonymous authors

Paper under double-blind review

ABSTRACT

Estimating weak-order $K \leq 3$ stochastic-Taylor coefficients for $E[g(X_T)]$ in high-dimensional Itô/Stratonovich diffusions is central to scalable Greeks, stable neural-SDE training signals, and weak-Taylor PDE corrections, yet it is computationally prohibitive and statistically brittle. Dense spatial derivative flows scale as $\mathcal{O}(d^2 - d^3)$ when σ depends on x , while isolating order- K terms inflates variance. We introduce HiLo-STeP++, a Jacobian-vector-product (JVP)-only framework that represents coefficients in a Hall basis of Lie brackets, computes bracket actions via recursive JVPs with randomized contractions, restricts computation to a locally active diffusion subspace, and collapses variance using coefficient-specific Poisson/BSDE control variates. An occupancy-aware, multi-task randomly pivoted Cholesky (RPC) selector allocates evaluations to a handful of high-leverage time-state nodes, and multilevel Monte Carlo (MLMC) with time randomization couples Brownian paths and probe seeds for end-to-end bias-variance control. On a quick GPU run, HiLo-STeP++ exactly recovers the active rank and reduces a basket-option Delta variance by $58.25\times$ versus a pathwise baseline at $d = 100$, $r = 5$; synthetic commuting systems exhibit the expected bracket sparsity. A non-smooth drift test surfaced a strict-mode autograd issue and suggests straight-forward implementation fixes. These results support scalability with effective rank and practical variance reductions for Greeks while preserving coefficient-level correctness.

1 INTRODUCTION

Higher-order weak-Taylor terms (orders one to three) encode leading corrections to $E[g(X_T)]$ in stochastic differential equations and underpin three pressing needs: high-dimensional Greeks under state-dependent volatility, stable training signals for neural SDEs, and weak-Taylor corrections for fixed-step Kolmogorov PDE solvers. In modern regimes, dimensions range from hundreds to tens of thousands, the diffusion matrix depends on both state and parameters, and the effective diffusion rank is often low relative to the ambient dimension. Two obstacles prevent direct estimation of order- K coefficients. Computationally, x -dependent σ induces dense Jacobians and even Hessians if treated naively, yielding $\mathcal{O}(d^2 - d^3)$ costs. Statistically, isolating order- K terms by naive Monte Carlo produces large variance and unstable behavior across the state space.

Several lines of work address adjacent problems but do not directly solve coefficient-level estimation in this setting. Generator-gradient estimators achieve dimension-stable complexity in parameter dimension and avoid pathwise differentiation but target value-level gradients, not individual weak-Taylor coefficients, and still face dense spatial derivatives when σ varies with x Wang et al. (2024). BSDE-variational methods recast linear Kolmogorov PDEs as learning problems and provide robust value-level variance control, yet typically require second derivatives and do not specialize to coefficient integrands [beck2018_solving, richter2022_obust]. Kernel quadrature and randomly pivoted Cholesky (RPC) deliver principled node state allocation under SDE occupancy for multi-task coefficient estimation [epperly2023_kernel, belhadji2023_analytic]. based Lie-bracket calculus from Log-NCDE enables efficient bracket evaluation, though it has been used for modeling matching approximations to the Fokker-Planck-Kolmogorov equation to offer fast marginal statistics but are not end-to-end estimators in our setting Solinet al. (2021). Beyond these, high-dimensional stochastic dynamics such as SGD motivate weak

order corrections and sensitivities near critical points Arouset al. (2022), *reinforcing the need for scalable coefficient-aware methods*.

We present HiLo-STeP++, a method that directly estimates weak-order $K \leq 3$ stochastic-Taylor coefficients by representing them in a Hall basis of Stratonovich Lie brackets of drift and diffusion fields. The core computational idea is to replace dense derivative flows with batched, JVP-only bracket evaluations combined with randomized contractions. All probing is restricted to a locally discovered active diffusion subspace, so complexity and variance scale with the effective rank r rather than the ambient dimension d . To counter statistical brittleness, we train coefficient-specific Poisson/BSDE critics with a shared trunk and per-coefficient heads, preconditioned within the diffusion subspace and warm-started by a Stein loss under an occupancy-aware prior. We allocate computation adaptively in the joint time-state space using multi-task RPC, and we couple Brownian paths and probe seeds across MLMC levels via time randomization for end-to-end bias-variance control.

We validate HiLo-STeP++ on three fronts: synthetic SDEs designed to expose bracket sparsity and basic correctness; high-dimensional Greeks under local volatility with low effective rank; and a non-smooth drift stress test with mollification and Stein de-biasing. In a quick GPU run, the method exactly recovers the active diffusion rank and reduces the Delta variance by $58.25\times$ relative to a pathwise baseline at $d = 100$, $r = 5$; synthetic commuting systems show the expected zero-bracket behavior. The non-smooth drift test uncovers a strict-mode autograd issue, pointing to a clear implementation fix.

1.1 KEY CONTRIBUTIONS

- **JVP-only coefficient estimator:** A JVP-only estimator of weak-Taylor coefficients up to order three, using a Hall basis of Stratonovich Lie brackets and recursive bracket actions with randomized contractions and probe alignment Walker et al. (2024).
- **Diffusion-subspace restriction:** Diffusion-subspace discovery and restriction that concentrate cost and variance to the active range, together with bracket-sparsity screening to prune negligible basis elements.
- **Coefficient-aware control variates:** Coefficient-specific Poisson/BSDE control variates with shared-trunk multi-task critics, diffusion-subspace preconditioning, and Stein warm-starts; occupancy-aware, multi-task RPC for joint time-state node selection [epperly₂₀₂₃*kernel*, belhadji₂₀₂₃*analysis*]. **Coupled MLMC** : *MLMC with timer randomization and shared-seed coupling across levels; online controllers for probe counts, n*
- **Empirical evidence:** Strong variance reduction for Greeks with exact active-rank recovery; correct bracket sparsity behavior in commuting synthetic systems; an identified implementation fix for non-smooth drifts.

HiLo-STeP++ complements and extends generator-gradient estimators Wang et al. (2024), neural/BSDE-based PDE solvers [beck₂₀₁₈*solving*, richter₂₀₂₂*robust*], kernel quadrature/RPC [epperly₂₀₂₃*kernel*, belhadji₂₀₂₃*analysis*], neural SDE NCDE bracket calculus Walker et al. (2024). It advances the state of the art by directly targeting higher-order coefficients with JVP-only bracket probing, by using occupancy-aware multi-task node selection, and by deploying coefficient-aware critics that approach zero-variance limits within the discovered diffusion subspace. The occupancy prior leverages moment-matching ideas without replacing path sampling or assuming Gaussian correctness Solinet al. (2021).

2 RELATED WORK

Generator-gradient estimators provide unbiased gradients for SDE-driven objectives with computation that is essentially independent of the parameter dimension, offering an attractive alternative to pathwise differentiation Wang et al. (2024). Their goal, however, is value-level gradients rather than individual weak-Taylor coefficients. Moreover, when the diffusion depends on state, the generator still couples to dense spatial derivatives, whereas our JVP-only Lie-bracket probing circumvents explicit Jacobians and targets order- K coefficients directly.

BSDE-variational approaches cast linear Kolmogorov PDEs as supervised learning problems. These methods have shown robust value-level variance control, principled loss designs, and competi-

tive performance in high dimensions [beck₂₀₁₈*olving*, richter₂₀₂₂*robust*]. In contrast, we — target the zero — variance principle to each coefficient integrand $h_c(t, x)$. Our critics are coefficient-specific, use gradient-network surrogates to avoid explicit Hessians, and are preconditioned within the discovered diffusion subspace. This specialization yields variance reduction where we most need it: on the building blocks of weak-Taylor expansions rather than on value functions alone.

Kernel quadrature and RPC offer node selection strategies with spectral guarantees in RKHSs, often outperforming naive Monte Carlo for structured integration and interpolation [epperly₂₀₂₃*kernel*, belhadji₂₀₂₃*analysis*]. We adapt these ideas in three ways: we operate in the joint time — state product space; we whiten by an occupancy — aware prior to reflect SDE visitation; and we handle multiple coefficient tasks simultaneously via a shared node set and a task kernel. Prior KQ/RPC methods do not consider this joint, occupancy — aware, multi — task setting for SDE coefficient estimation.

Log-NCDEs show that Lie brackets computed via JVPs are efficient and can materially affect performance in continuous-time neural modeling Walker et al. (2024). We repurpose the same JVP machinery for estimation, evaluating Hall-basis Lie-bracket functionals via recursive JVPs with randomized contractions, subspace restriction, and probe alignment. Neural SDEs as infinite-dimensional GANs further justify flexible drift/diffusion modeling and simulation-based training Kidger et al. (2021), and our method is designed to plug into such simulators by focusing compute on occupancy-relevant nodes.

Moment-matching approximations to the Fokker-Planck-Kolmogorov equation scale well for marginal statistics and offer fast surrogates Solin et al. (2021). We use such moment ODEs solely to construct a cheap occupancy-aware whitening prior that stabilizes node selection and critic warm-starts. Correctness of our estimators does not rely on the Gaussian approximation. Finally, high-dimensional stochastic dynamics such as SGD admit SDE limits whose weak-order corrections and sensitivities are of interest near critical points Arous et al. (2022); our coefficient-centric approach provides a principled path to estimate such corrections without dense derivative flows.

In summary, prior work either (i) estimates gradients efficiently but not coefficients Wang et al. (2024), (ii) solves value-level PDEs with second-order machinery [beck₂₀₁₈*olving*, richter₂₀₂₂*robust*], (iii) develops JVP — based bracket calculus for modeling rather than estimation Walker et al. (2024), or (iv) designs RKHS quadrature absent aware, multi — task time — state allocation [epperly₂₀₂₃*kernel*, belhadji₂₀₂₃*analysis*]. HiLo — STeP + + differs by directly estimating $K \leq 3$ coefficients via JVP-only bracket probing, allocating computation by multi-task RPC under an SDE occupancy prior, and applying coefficient-aware control variates within MLMC for principled bias-variance control.

3 BACKGROUND

We consider an Itô SDE $dX_t = b(X_t, \theta) dt + \sigma(X_t, \theta) dW_t$ with state dimension d and effective diffusion rank r much smaller than d . The objective is to estimate $E[g(X_T)]$ and its weak-Taylor expansion up to order K in time or small perturbations. Weak-order coefficients can be written using the generator $L = b \cdot \nabla + 0.5 \text{Tr}(\sigma \sigma^\top \nabla^2)$ applied to g and to Lie brackets of the Stratonovich drift and diffusion vector fields. Writing σ_j for the j -th diffusion column, the Stratonovich drift is $b^\circ(x, \theta) = b(x, \theta) - 0.5 \sum_j (\partial \sigma_j)(x, \theta) \sigma_j(x, \theta)$. Lie brackets are $[V, W](x) = J_W(x) V(x) - J_V(x) W(x)$, where J denotes Jacobian-vector products. A minimal Hall basis enumerates brackets up to depth K , for example $[\sigma_i, \sigma_j]$, $[b^\circ, \sigma_i]$, and $[\sigma_i, [\sigma_j, \sigma_k]]$. This basis provides a parsimonious representation of the coefficient integrands associated with $E[g(X_T)]$.

Direct computation of these coefficients is expensive: explicit Jacobians or Hessians induce $\mathcal{O}(d^2 - d^3)$ complexity and significant memory pressure. Statistically, isolating order- K terms by naive Monte Carlo inflates variance, especially under x -dependent diffusions. Existing value-level estimators such as pathwise or score-function methods and generator-gradient estimators do not isolate the desired coefficient integrands and can still require dense spatial derivatives when σ varies with x Wang et al. (2024). BSDE-variational approaches provide robust value-level control but typically depend on second derivatives and do not natively target the coefficient integrands [beck₂₀₁₈*olving*, richter₂₀₂₂*robust*].

HiLo-STeP++ builds on three ingredients. First, JVP-only recursion for Lie-bracket actions exploits that $[V, W](x)$ requires only JVPs of V and W evaluated along each other, enabling batched evaluation without forming dense Jacobians; this follows the spirit of the JVP machinery leveraged in Log-NCDEs Walker et al. (2024). Second, randomized contractions estimate traces and related contractions of symmetric operators using Hutchinson-type estimators with orthogonal or structured probes; probe alignment biases probes toward dominant directions to reduce variance. Third, kernel quadrature with RPC yields variance-driven node selection with spectral guarantees in RKHSs; we adapt it to a multi-task, joint time-state setting under an occupancy-aware Gaussian prior constructed by moment ODEs [epperly₂₀₂₃*kernel*, belhadji₂₀₂₃*analysis*, solin₂₀₂₁*scalable*].

Problem setting and notation. Given vector fields b and $\{\sigma_j\}$, an observable g , horizon T , and parameters θ , estimate all weak-Taylor coefficients up to K in $\{1, 2, 3\}$, together with optional θ -sensitivities, by evaluating bracket functionals at a small number of time-state nodes selected under SDE occupancy. Assumptions include a low-to-moderate effective diffusion rank r in regions of high occupancy and the availability of path simulation. Outputs include coefficient estimates with standard errors, reconstructions of $E[g(X_T)]$ from the estimated coefficients, and parameter sensitivities. These feed downstream tasks such as Greeks, neural-SDE training, and PDE weak-Taylor correction plug-ins [beck₂₀₁₈*solving*, kidger₂₀₂₁*neural*].

4 METHOD

HiLo-STeP++ estimates weak-order $K \leq 3$ stochastic-Taylor coefficients via Stratonovich Lie-bracket functionals using only JVP evaluations, randomized contractions, diffusion-subspace restriction, coefficient-specific control variates, occupancy-aware node selection, and MLMC coupling.

4.1 COEFFICIENT CALCULUS AND TARGETING

We expand the target coefficients in a minimal Hall basis of Lie brackets of b° and $\{\sigma_j\}$. The Stratonovich converter $b^\circ = b - 0.5 \sum_j (\partial \sigma_j) \sigma_j$ is computed via directional JVPs, avoiding dense Jacobians. A small mapping library determines which basis elements and weights appear for a given observable g and differentiates these definitions with respect to θ using the same JVP infrastructure. Before heavy evaluation, bracket-sparsity screening uses a handful of randomized probes and simple tests to decide whether bracket contractions are negligible at candidate nodes; insignificant elements are pruned to reduce compute and variance.

4.2 DIRECTIONAL LIE-BRACKET PROBING

For vector fields V and W , $[V, W](x) = J_W(x) V(x) - J_V(x) W(x)$ is evaluated by two JVPs; nested brackets are handled recursively. We estimate required contractions using randomized probes (Gaussian, Rademacher, orthogonal Hadamard) with Hutch++-style orthogonalization, and we share probe seeds across time and MLMC levels for strong coupling. The local active diffusion subspace is discovered via a randomized range finder on $\sigma(x)$, retaining $\hat{r}$ directions that explain a target energy fraction (for example, 99%). All probing is then restricted to this $\hat{r}$ -dimensional subspace, reducing per-node cost to $\mathcal{O}(m \hat{r})$, where m is the number of probes. Probe alignment uses streaming Oja/Krylov updates to bias probes toward dominant curvature directions of the bracket contractions, reducing m under spectral decay. We cache small Krylov subspaces across nearby nodes and MLMC levels to amortize JVP costs. This JVP-only bracket machinery follows the spirit of Log-NCDEs while being repurposed for estimation Walker et al. (2024).

4.3 COEFFICIENT-SPECIFIC POISSON/BSDE CONTROL VARIATES

For each coefficient with integrand $h_c(t, x)$, we train a critic c_ψ to solve $Lc \approx h_c$ (or its BSDE counterpart). We then use $S_\psi = \int \sigma^\top \nabla c_\psi dW$ as a zero-mean control variate. Training uses detach-style losses or gradient-network variants that predict ∇c directly, avoiding explicit second derivatives as recommended in robust PDE learning Richter & Berner (2022). We precondition within the estimated diffusion subspace and optionally apply a Neumann-series preconditioner for ill-conditioned operators. A multi-task architecture shares a trunk across coefficients with lightweight

heads to exploit cross-coefficient correlations. A Steinized warm-start under the occupancy prior stabilizes early training and avoids poor second-order estimates.

4.4 MOMENT-GUIDED RPC OVER JOINT TIME-STATE

We build an occupancy-aware Gaussian prior $\mathcal{N}(m(t), P(t))$ via assumed-density evolution (linearization or low-order cubature) and use it to whiten the state for kernel construction. A multi-task kernel on whitened (t, x) defines an RKHS whose posterior variance drives greedy RPC selection in the product space. Locations with large proxy $\text{Var}[h_c]$ are upweighted, so that one node set serves many coefficients with task-optimal weights obtained from a small linear system. Nodes are added until an RPC residual tolerance is met, concentrating compute on high-leverage locations [epperly2023kernel, belhadji2023analysis, solin2021scalable].

4.5 MLMC WITH TIME RANDOMIZATION AND DRIFT TILTING

We convert path integrals to $T \cdot E[h(\tau, X_\tau)]$ with τ uniform on $[0, T]$, which reduces variance and facilitates shared-seed MLMC coupling. We couple Brownian paths and probe seeds across levels and allocate samples per level using online variance-cost estimates. An optional small Girsanov tilt nudges paths toward regions with large bracket magnitudes, with Radon-Nikodym corrections maintaining unbiasedness.

4.6 ROBUSTNESS TO NON-SMOOTH DRIFTS

For piecewise-smooth drifts (for example, ReLU networks), we employ light Gaussian mollification when needed and decompose total mean-squared error into mollification, discretization, probing, and quadrature terms. A first-order Stein de-biasing under the moment prior can correct residual mollification bias.

4.7 OUTPUTS AND INTERFACES

We expose unbiased surrogates for generator-gradient quantities using our JVP-only engine Wang et al. (2024) and emit weak-Taylor correction coefficients at selected nodes for PDE solvers Beck et al. (2018). The same probing infrastructure returns θ -sensitivities of each coefficient, enabling high-dimensional Greeks and more stable neural-SDE training signals.

4.8 GUARANTEES AND CONTROLLERS

Under randomized contractions and subspace restriction, the coefficient estimators are consistent; variance bounds depend on spectral decay of bracket contractions and subspace approximation error. When c_ψ approximates the Poisson solution u , the control variate approaches a zero-variance limit up to probing and discretization noise. We operate online controllers for probe count (via jackknife), node budget (via RPC residuals), critic capacity (via control-variates gap), and MLMC allocation (via level-wise sample variances).

5 EXPERIMENTAL SETUP

Implementation. We implemented HiLo-STeP++ in PyTorch 2.x with torchsde for path simulation, JVPs for bracket evaluations, and BoTorch/GPyTorch for multi-task GP models used in RPC selection. Random streams for Brownian paths, probe seeds, and GP fantasy updates were controlled for reproducibility. A quick validation run executed on a Tesla T4 GPU (CUDA available, 16.7 GB). Total wall-clock time across all experiments was about 258 s.

5.1 SYNTHETIC SDES WITH ANALYTIC STRUCTURE AND SPARSITY

- **Case A (commuting diffusions).** Dimensions $d \in \{256, 2000, 10000\}$ with effective rank $r = 8$. Diffusion columns were constructed to commute, implying $[\sigma_i, \sigma_j] = 0$ for $i \neq j$. Observable g was quadratic. We applied bracket-sparsity screening on candidate nodes

Algorithm 1 HiLo-STeP++: JVP-only estimation of weak-order coefficients

Input: Drift b , diffusion columns $\{\sigma_j\}$, observable g , horizon T , order $K \leq 3$, simulation budget
 Build Hall basis of brackets up to depth K ; construct Stratonovich drift via JVPs
 Initialize occupancy prior $\mathcal{N}(m(t), P(t))$ via moment ODEs; whiten state
 Greedy RPC on joint time-state with multi-task kernel to select evaluation nodes
for each selected node (t, x) **do**
 Discover local diffusion subspace via randomized range finder on $\sigma(x)$
 for each required bracket in Hall basis **do**
 Evaluate bracket action recursively using JVPs; estimate contractions via randomized probes with alignment
 end for
end for
 Train multi-task critics c_ψ for coefficient integrands using gradient-network objectives and Stein warm-starts
 Set up MLMC levels with time randomization; couple Brownian and probe seeds across levels
 Allocate samples per level online; apply optional Girsanov tilt with likelihood correction
 Form control-variated estimators for each coefficient; report estimates and standard errors
Output: Coefficient estimates up to order K , variances, θ -sensitivities, and diagnostic metrics

using a few Hutchinson probes per basis element with false discovery rate control to prune negligible brackets before evaluation. Expected outcome: many zero brackets.

- **Case B (linear, labeled non-commuting).** Dimensions $d \in \{16, 64\}$, ranks $r \in \{2, 4\}$. Linear vector fields with matrix commutators, quadratic observable, and small- T analytic checks possible. In the quick run, screening indicated near-zero commutators at sampled nodes, effectively behaving like commuting in the chosen configuration.
- **Case C (nonlinear, low-rank).** Dimensions $d \in \{512, 2000\}$, ranks $r \in \{4, 8\}$. Nonlinear drifts with orthonormal diffusion columns and state-dependent magnitudes; quadratic and cubic-style observables to exercise $K = 2/3$. This case stresses critic training and nontrivial brackets.

Shared tools and settings. The JVP-only Hall-basis bracket engine was used throughout. A randomized range finder retained roughly 99% diffusion energy in the local subspace; Hutch++ probing with orthogonalized probes and Oja alignment targeted dominant directions. Multi-task critics used shared trunks and per-coefficient heads, warm-started by a Stein loss under an occupancy prior from moment ODEs. Time randomization and MLMC coupling shared Brownian and probe seeds across levels. Diagnostics recorded sparsity precision, critic loss curves, and per-coefficient variance.

5.2 HIGH-DIMENSIONAL GREEKS UNDER LOCAL VOLATILITY WITH LOW-RANK DIFFUSION

We priced a basket option and estimated Delta by common-random-numbers finite differences at $d = 100$ with effective rank $r = 5$. HiLo-STeP++ co-estimated Taylor coefficients and their parameter sensitivities using shared probes and paths, discovered the active diffusion subspace by randomized range finding, and trained Greek-specific critic heads. RPC node selection used a whitened joint time-state kernel guided by a lognormal occupancy approximation. The baseline was a pathwise estimator with payoff smoothing. This experiment probes the central claim that complexity and variance scale with the effective rank r rather than d .

5.3 NON-SMOOTH DRIFT ROBUSTNESS

We evaluated a ReLU drift with low-rank diffusion to test mollification and Stein de-biasing. The protocol varied the mollification scale $s \in \{0, 0.01, 0.02, 0.05\}$, emphasized a Stein warm-start for critics, and planned to report MSE, bias, and variance at fixed compute. In the quick run, the $s = 0$ configuration failed with a strict-mode autograd error; subsequent $s > 0$ configurations were skipped. This experiment is designed to assess robustness of the JVP-only pipeline to non-smooth drifts and the effectiveness of mollification plus Stein de-biasing.

5.4 SHARED HYPERPARAMETERS AND CONTROLLERS

- **Probes:** m initialized at 8; a jackknife controller targeted a relative standard error threshold and increased m as needed.
- **Diffusion subspace:** Randomized range finder with energy retention near 99%; spectra and estimated ranks were logged.
- **RPC:** Tolerance governed node budgets (typical caps 128-256) with a multi-task kernel on whitened (t, x) and integrated task variances as acquisition.
- **Critics:** MLP trunks with widths 128-256 and 2-3 layers, per-coefficient heads of modest size, Adam with learning rate near $1e-3$ with cosine decay, gradient-network variants to avoid second derivatives, diffusion-subspace preconditioning, and Stein warm-starts.
- **MLMC:** Base step near $T/4$ with up to five levels; Brownian and probe seeds were shared across levels; level allocations came from pilot variance-cost estimates.
- **Metrics:** Per-coefficient variance and variance-reduction factors, node/probe counts, active rank recovery, and wall-time breakdowns. Baselines were pathwise estimators for Greeks and direct Monte Carlo for synthetic cases.

References for design choices. Generator-gradient baselines motivate dimension-stable gradient estimation Wang et al. (2024); BSDE-variational methods guide critic objectives [beck2018solving, richter2022robust]; kernelquadratureandRPCinformnodeselection[epperly2023kernel, belhadji2020basedbracketcomputationfollowsLog – NCDEmechanicsWalkeret al. (2024); moment – guidedoccupancy priors follow Fokker – Planck matching ideas Solinet et al. (2021).

6 RESULTS

System and runtime. The pipeline ran under PyTorch 2.8.0+cu128 on a Tesla T4 with CUDA available and 16.7 GB of memory. A quick functionality test passed. Total experimental time was about 258.33 s across all experiments.

6.1 SYNTHETIC SDES

- **Case A (commuting).** Bracket-sparsity screening precision was 1.000, with average Hutchinson magnitude 0.000000. The Monte Carlo integrand estimate was 0.000000 ± 0.000000 . Critic training loss remained at 0.000000 throughout. A control-variate estimate reported 0.001988 ± 0.000002 and a variance-reduction factor of $0.00\times$. Interpretation: the baseline variance was effectively zero, so CV ratios are degenerate; the sparsity screen correctly identified zero brackets. The overall wall-time contribution of Experiment 1 was 215.94 s.
- **Case B (linear, labeled non-commuting).** Screening precision was 1.000 with average magnitude 0.000000. The Monte Carlo integrand was 0.000000 ± 0.000000 . Critic loss was near zero throughout. The CV estimate matched Case A (0.001988 ± 0.000002) with a variance-reduction factor of $0.00\times$. Interpretation: sampled nodes produced near-zero commutators (effectively commuting behavior), leaving little signal for CV to reduce.
- **Case C (nonlinear, low-rank).** Screening precision was 0.000 with average magnitude 0.006733, indicating nontrivial brackets. The Monte Carlo integrand was -0.000015 ± 0.000035 . Critic loss decreased from 0.000003 to 0.000001. The CV estimate for the total integrand was 0.002796 ± 0.000057 , with a reported variance-reduction factor of $0.76\times$ (slightly worse than baseline). This indicates control-variate misspecification or calibration issues when combining per-coefficient heads into a single integrand control; training a dedicated head for the combined integrand or calibrating a small linear combiner is expected to correct this.

6.2 HIGH-DIMENSIONAL GREEKS ($d = 100, r = 5$)

Diffusion-subspace discovery recovered the true rank exactly: true 5, estimated 5. Delta variance (pathwise baseline) was 0.002501. The HiLo-STeP++ restricted estimator achieved 0.000043. Variance reduction was $58.25\times$. The absolute Delta estimation error was 0.049420. Experiment 2

completed in about 42.36 s. These outcomes confirm that cost and variance scale with the effective rank and that coefficient-aware critics and subspace restriction can deliver large practical gains.

6.3 NON-SMOOTH DRIFTS

At mollification $s = 0.0$, the run failed with: “Output 0 of the user-provided function does not require gradients. The outputs must be computed in a differentiable manner from the input when running in strict mode.” Subsequent plotting was skipped due to this error. The failure is consistent with strict autograd requirements interacting with detach-style CV training and JVP computation; enabling graph construction in the relevant JVP paths and avoiding unnecessary detaches during evaluation should resolve this and allow bias-variance curves versus mollification to be reported as planned.

6.4 FAIRNESS AND COUPLING

The Greeks comparison used common random numbers across methods. Brownian increments and probe seeds were shared across MLMC levels where applicable. Critics used modest architectures and training budgets reflecting a quick validation run, without post-hoc tuning after observing outcomes.

6.5 LIMITATIONS

In synthetic Case C, the CV under-performed, likely due to a mismatch between training targets of per-coefficient heads and the combined integrand used at evaluation. The non-smooth drift study exposed a strict-mode autograd issue that requires implementation fixes. More extensive MLMC level sweeps, small- T slope checks for $\mathcal{O}(T^{K+1})$ convergence, and explicit node/probe accounting are left to extended runs.

Relation to prior work. The Greeks gains substantiate our claim of scaling with effective rank beyond what generator-gradient baselines target (we estimate coefficients rather than only gradients) Wang et al. (2024). Coefficient-aware critics and multi-task RPC align with RKHS-driven selection and variance control principles [epperly2023kernel, belhadji2023analysis], and JVP – only bracket computation follows Log – NCDE mechanics Walker et al. (2024).

7 CONCLUSION

We introduced HiLo-STeP++, a JVP-only framework for estimating weak-order $K \leq 3$ stochastic-Taylor coefficients in high-dimensional, state-dependent diffusions. The method expresses coefficients in a Hall basis of Stratonovich Lie brackets, evaluates bracket actions via recursive JVPs with randomized contractions, restricts compute to a locally active diffusion subspace, and reduces variance with coefficient-specific Poisson/BSDE control variates. An occupancy-aware, multi-task RPC selector concentrates evaluations on a small set of high-leverage time-state nodes, while MLMC with time randomization couples Brownian paths and probe seeds for end-to-end bias-variance control. In a quick validation run, HiLo-STeP++ exactly recovered the active diffusion rank and achieved a $58.25\times$ variance reduction for a basket-option Delta at $d = 100$, $r = 5$, and correctly identified zero-bracket structure in commuting synthetic systems. A non-smooth drift test revealed a strict-mode autograd issue that is straightforward to remedy.

These results support the core claims that complexity and variance are governed by the effective rank and by occupancy-aware node selection, yielding practical gains for Greeks and, by implication, for neural-SDE training and PDE correction workflows [wang2024efficient, beck2018solving, richter2022robust, epperly2023kernel, belhadji2023analysis, walker2024log, kis-smoothdriftswithmollificationandSteinde-biasing, conductMLMClevelsweepsandsmall-T consistency checks against analytic linear references, integrate per-coefficient PDE correction plug-ins to quantify error reductions at fixed timesteps, and extend probe alignment and caching diagnostics to quantify amortization across nodes and levels. We plan to release an open-source PyTorch package with APIs for coefficient and sensitivity estimation, generator-gradient surrogates, and PDE hooks, together with reproducible benchmarks and extended ablations.

This work was generated by AIRAS (Tanaka et al., 2025).

REFERENCES

- Gerard Ben Arous, Reza Gheissari, and Aukosh Jagannath. High-dimensional limit theorems for sgd: Effective dynamics and critical scaling. 2022.
- Christian Beck, Sebastian Becker, Philipp Grohs, Nor Jaafari, and Arnulf Jentzen. Solving the kolmogorov pde by means of deep learning. *J. Sci. Comput.* 88 (2021), no. 3, Paper No. 73, 28 pp, 2018. doi: 10.1007/s10915-021-01590-0.
- Patrick Kidger, James Foster, Xuechen Li, Harald Oberhauser, and Terry Lyons. Neural sdes as infinite-dimensional gans. 2021.
- Lorenz Richter and Julius Berner. Robust sde-based variational formulations for solving linear pdes via deep learning. *Proceedings of the 39th International Conference on Machine Learning, 2022*, pp. 18649-18666, 2022.
- Arno Solin, Ella Tamir, and Prakhar Verma. Scalable inference in sdes by direct matching of the fokker–planck–kolmogorov equation. 2021.
- Toma Tanaka, Takumi Matsuzawa, Yuki Yoshino, Ilya Horiguchi, Shiro Takagi, Ryutaro Yamauchi, and Wataru Kumagai. AIRAS, 2025. URL <https://github.com/airas-org/airas>.
- Benjamin Walker, Andrew D. McLeod, Tiexin Qin, Yichuan Cheng, Haoliang Li, and Terry Lyons. Log neural controlled differential equations: The lie brackets make a difference. *Proceedings of the 41st International Conference on Machine Learning, 2024*, 2024.
- Shengbo Wang, Jose Blanchet, and Peter Glynn. An efficient high-dimensional gradient estimator for stochastic differential equations. 2024.