TAHES++: TANGENT-AWARE HIGHER-ORDER ELLIPTICAL SMOOTHING WITH PROVABLE, SPOOF-RESISTANT SUBSPACE CONTROL

Anonymous authors
Paper under double-blind review

ABSTRACT

We study certified robustness at scale under off-manifold attacks that spoof randomized smoothing. Classical smoothing provides scalable L_2 guarantees but suffers dimension-driven information limits and ignores data geometry, while recent first-order certificates largely assume isotropy and remain blind to normal directions exploited by Shadow-style attacks. We propose TAHES++, a tangent-aware, anisotropic first-order certification framework that preserves randomized smoothing validity while aligning the smoothing covariance with the data manifold. We learn a global tangent-normal split, form a fixed ellipsoidal covariance, whiten to isotropic space, extend first-order certification with projected (tangent/normal) constraints, and clamp normal sensitivity via scalable Lipschitz upper bounds. We adapt transitivity to the ellipsoidal metric and introduce a light training regularizer that flattens the smoothed classifier off-manifold. Theoretically, whitening reduces to the isotropic case; subspace constraints and Lipschitz clamping are conservative, and directional control ties robustness to the manifold dimension [cohen$_2$019$_c$$ertified, yang_2020_randomized, mohapatra_2020_higher, kumar_2020_curse, xu_2024_eclipse].Empiric$
$to-endandexecutedasyntheticdryrunconfirmingfeasibility, stablegeometryextraction, andhighcleanaccur$
$10 and ImageNet - 100 against RS, RS4A, HO -$
$Cert, and GICR, reportingellipsoidalcertificates, L_2$ summaries, and a spoofability index with confidence intervals [cohen$_2$019$_c$$ertified, yang_2020_randomized, mohapatra_2020_higher, cullen_2022_double, ghiasi_2020_breaking]$.

1 INTRODUCTION

Randomized smoothing converts a base classifier into a smoothed one with provable L_2 robustness and remains the only practically scalable certified defense on large vision models [cohen$_2$019$_c$$ertified].However, whenrestrictedtolabelstatistics, certifiedradiidegradeunfavorablywithdimension,$
$innaturalimages, semanticvariationisoftenconcentratedalonglow -$
$dimensionaltangentsofadatamanifold, whilenuisanceorunnaturalvariationliespredominantlyalongnormaldirecti$
$ordercertificationaugmentssmoothingbyleveraginggradientinformationtoenlargecertifiedsets, buttypicalformul$
$styleoff-manifoldattacksshowthatmodelscanbemanipulatedtobothmisclassifyandreturnspoofedcertificates, exp$
$andLipschitz-basedcertificationpromisesstrongercontrol, butoftenreliesonheavysecond-$
$ordermachineryorsemidefiniteprogramsthatimpedescale[singla_2020_second, chen_2020_semialgebraic, latorre_2020_{li}$
space and does not incorporate manifold structure [cullen$_2$022$_d$$ouble]$.

We introduce TAHES++, a tangent-aware, anisotropic first-order certification framework with provable, spoof-resistant subspace control. The central idea is to preserve randomized smoothing validity by fixing the smoothing distribution, but learn an input-independent ellipsoidal covariance aligned to the data manifold. After whitening to isotropic space, we perform first-order certification with explicit tangent and normal constraints, clamp normal sensitivity via conservative Lipschitz upper bounds, and extend transitivity in the same ellipsoidal metric. We also propose a light training regularizer that reduces off-manifold gradients while preserving tangent responsiveness. This combination is designed to enlarge certificates along semantic directions, maintain conservatism off manifold, and resist spoofing.

Concretely, we: (1) estimate a global tangent basis by aggregating local decoder-Jacobian tangents from a frozen generative backbone, (2) construct a fixed anisotropic covariance and certify in the induced Mahalanobis geometry via whitening, (3) extend first-order certification to use projected gradient constraints in tangent and normal subspaces, (4) clamp normal sensitivity using scalable Lipschitz upper bounds that remain conservative, (5) adapt transitivity to search along learned tangent axes in whitened space and union ellipsoidal certificates, and (6) introduce a small training penalty that flattens the smoothed classifier off manifold while keeping the smoothing distribution fixed. Theoretical validity follows from maintaining a fixed smoothing distribution, equivalence via whitening, and conservative confidence bounds; direction-wise control ties robustness to manifold dimension rather than ambient dimension [cohen$_{2}019_{c}ertified, yang_{2}020_{r}andomized, mohapatra_{2}020_{h}igher, kumar_{2}020_{c}urse$].

We verify feasibility through an end-to-end synthetic dry run and lay out a comprehensive evaluation on CIFAR-10 and ImageNet-100 against RS, RS4A, HO-Cert, and HO-Cert+GICR [cohen$_{2}019_{c}ertified, yang_{2}020_{r}andomized, mohapatra_{2}020_{h}igher, cullen_{2}022_{d}ouble].Wereportellipsoidalcertificate$ summaries, and a spoofability index to quantify directional sensitivity and calibration, especially under normal-subspace attacks inspired by Shadow [ghiasi$_{2}020_{b}reaking$].

Our contributions are:

- **Anisotropic smoothing aligned to geometry.** Tangent-aware anisotropic smoothing with a fixed, input-independent covariance aligned to data geometry, preserving classical randomized smoothing validity [cohen$_{2}019_{c}ertified, yang_{2}020_{r}andomized$].**First-order in the Mahalanobis metric.**$First-ordercertificationintheMahalanobismetricwithprojectedtangent/normalconstraints, yieldingdirectional$

- **Provable normal clamping.** Provable normal clamping via scalable Lipschitz upper bounds to mitigate Shadow-style off-manifold inflation while remaining practical [xu$_{2}024_{e}clipse, latorre_{2}020_{l}ipschitz, ghiasi_{2}020_{b}reaking$].**Geometry-aware transitivity and regularizer.**$Geometry-awaretransitivityalonglearnedtangentaxesandalighttrainingregularizer(TaReS)thatflattensoff-manifoldsensitivitywithoutchangingthesmoothingdistribution[cullen_{2}022_{d}ouble, vaishnavi_{2}022_{a}ccelerating]$.

- **Transparent reporting.** A reporting protocol that outputs full ellipsoidal certificates, scalar L_2 summaries for comparability, and a spoofability index with confidence intervals to support transparent auditing.

Looking forward, TAHES++ seeks robust, transparent, and scalable certification aligned with semantics. Future work includes multi-subspace and perceptual geometries, tighter but scalable Lipschitz clamping, and integration with certified training and transfer frameworks [yang$_{2}020_{r}andomized, vaishnavi_{2}022_{a}ccelerating$].

2 Related Work

2.1 Randomized smoothing and its limits

Classical randomized smoothing certifies L_2 robustness under Gaussian noise and scales to large models [cohen$_{2}019_{c}ertified].RS4Ageneralizessmoothingviaaunifyingframeworkacrossnormsandcharacterizesoptima$ $basedlimitsforlabel-onlystatisticsthatdegradewithdimension[yang_{2}020_{r}andomized, kumar_{2}020_{c}urse].TAHES+$ $+preservessmoothingvaliditybutaugmentsitwithlearnedanisotropyandfirst-orderdirectionalinformationtomovebeyondlabel-onlylimitsinageometry-awaremanner.$

2.2 First-order certification

HO-Cert enlarges certified regions by estimating class score gaps and gradient norms with low variance, keeping the smoothing distribution fixed [mohapatra$_{2}020_{h}igher].However, itpresumesisotropyandlacksexplicitcontrolovertangentversusnormalsensitivity.T$ $+bringsfirst-ordercertificationintoananisotropicMahalanobismetricbywhiteningwithafixedcovarianceandintro$ $projectedconstraints, enablingdirectionaltighteningwithoutchangingtheunderlyingnoisedistribution.$

2.3 Curvature and Lipschitz-based certification

Curvature-based methods (e.g., CRC) derive certificates from Hessian bounds and can regularize training, yet require managing curvature via convex optimization or specialized machinery that limits scalability [singla$_2$020$_s$econd]$.Lipschitz - basedmethodsprovideglobalsensitivityupperbounds; semialgebraicorSDP - basedapproachescanbetighterbutcomputationallyexpensive, whereascompositionalestimateslikeECLipsEofferscа$ $+usesLipschitzestimatesconservativelytoclampnormalsensitivitywithinafirst - orderprogram, avoidingheavyoptimizationwhileimprovingspoofresistance.$

2.4 Transitivity

GICR demonstrates that uniting local certificates can improve robustness in pixel L_2 space [cullen$_2$022$_d$ouble]$.Weextendthisideatotheellipsoidalmetricinducedbythelearnedcovariance : inwhitenedspace, wesearchalongtangentaxesandunionEuclideanballs(mappingbacktoellipsoids), whichbetteralign$ $efficient.$

2.5 Attacks and spoofing

Shadow-style attacks show that certified defenses can be coerced into issuing misleading certificates under large, semantically plausible off-manifold perturbations [ghiasi$_2$020$_b$reaking]$.Thismotivatesexplicitmeasurementandcontrolofnormalsensitivity.Broaderevidencethatadve$ $normtransformationscanbreakmodels, furtherunderscorestheneedforprincipled, provabledefenses[carlini_2017_ad$

2.6 Certified training

IBP offers scalable verified training with appropriate schedules [gowal$_2$018$_t$he]$.Certifiedrobustnesstransferacceleratesrobusttrainingbyknowledgetransfer[vaishnavi_2022_acceler$ $itspecificallytargetsoff-manifoldsensitivityofthesmoothedclassifierwhilemaintainingcompatibilitywithRS-$ $styletraining.$

2.7 Positioning

Relative to RS4A and HO-Cert, TAHES++ aligns smoothing to manifold geometry and certifies in a learned Mahalanobis metric with explicit subspace constraints. Unlike curvature or Lipschitz certificates that can be heavy, TAHES++ uses fast Lipschitz bounds only to clamp normal sensitivity. Compared to GICR, our transitivity operates along learned tangent axes in the intrinsic geometry, promising higher sample efficiency and semantic alignment [cohen$_2$019$_c$$ertified, yang_2020_randomized, mohapatra_2020_higher, cullen_2022_double$].

3 Background

3.1 Problem setting and notation

Let f be a base classifier and g its smoothed counterpart obtained by adding noise to the input. Classical randomized smoothing with an isotropic Gaussian certifies an L_2 ball around an input x when the estimated score gap and noise level guarantee prediction invariance within a radius [cohen$_2$019$_c$ertified]$.Withonlylabelstatistics, thelargestcertifiableradiussuffersanunfavorabledependenceonthea$ $ordercertificationstrengthensguaranteesbyestimatinggradientinformationofthesmoothedclassifierandyieldstig$

3.2 Mahalanobis geometry and whitening

We introduce a fixed, global anisotropic covariance Σ aligned to a learned tangent-normal decomposition. Write $\Sigma = \sigma_T^2 B_T B_T^T + \sigma_N^2 B_N B_N^T$, where columns of B_T span a tangent subspace and B_N its orthogonal complement. Whitening via a matrix L with $LL^T = \Sigma$ maps inputs to $y = L^{-1}x$. In y-space the noise is standard normal, so certification reduces to the isotropic case; the certified region in x-space is an ellipsoid $\{\delta : \|L^{-1}\delta\|_2 \leq r_y\}$. For comparability with L_2-based literature, we also report a scalar summary $r\_L2 = r\_y \cdot \sqrt{\lambda\_\min(\Sigma)}$.

3.3 Projected sensitivities

For any vector v, define $P_T(v) = B_T B_T^T v$ and $P_N(v) = v - P_T(v)$. In y-space, we estimate the class score gap $g\_\mathrm{gap}(y)$ and the norms $\|P_T \nabla g(y)\|\_2$ and $\|P_N \nabla g(y)\|\_2$ using Monte Carlo with antithetic noise and control variates. High-confidence one-sided bounds on these quantities feed a first-order program to yield a conservative radius [mohapatra$_2$020$_h igher$].

3.4 Normal clamping via Lipschitz bounds

To increase spoof resistance, we cap normal sensitivity with an input-independent Lipschitz upper bound $\hat{L}$ from scalable estimators such as compositional ECLipsE or per-layer spectral norm products. Replacing the Monte Carlo upper bound on $\|P_N \nabla g\|$ with $\min(\text{CI\_upper}, \hat{L})$ is conservative and preserves validity [xu$_2$024$_e clipse, latorre_2 020_l ipschitz$].

3.5 Transitivity and unions

Certificates can be expanded by uniting local certifications at auxiliary points near x. In y-space, we search along learned tangent axes for auxiliary centers with sufficient confidence and certify each; the union of Euclidean balls in y maps to a union of ellipsoids in x [cullen$_2$022$_d ouble$].

3.6 Assumptions and scope

We assume semantics concentrate in an m-dimensional tangent subspace with bounded curvature, making anisotropic smoothing and directional certification meaningful. The covariance Σ is fixed and input independent, ensuring randomized smoothing validity [cohen$_2$019$_c ertified].Thegoalistoexposeandcontroldirectionalsensitivity, particularlyalongnormals, toimprovecali$ $stylespoofing[ghiasi_2 020_b reaking]$.

4 Method

4.1 Geometry learning

We estimate a global tangent basis B_T by aggregating local tangents derived from a frozen generative backbone. For each training example x with latent z, we compute the decoder Jacobian $J = \partial G / \partial z$ and extract top-k left singular vectors U_x of J as local tangents. Aggregating projectors $P\_x = U\_x U\_x^T$ across the dataset via streaming subspace iteration yields the dominant eigenspace of the mean projector $\bar{P}$; its columns define B_T. The normal space is the orthogonal complement spanned by B_N. Because this split is global and input independent, randomized smoothing validity is preserved [cohen$_2$019$_c ertified$].

4.2 Anisotropic smoothing and whitening

We define $\Sigma = \sigma_T^2 B_T B_T^T + \sigma_N^2 B_N B_N^T$ with anisotropy $\beta = \sigma_N / \sigma_T$ in a moderate range. Construct L with $LL^T = \Sigma$ and whiten inputs as $y = L^{-1} x$. In y-space, noise is standard normal, so first-order certification applies unchanged; certified Euclidean balls map back to ellipsoids in x.

4.3 First-order certification with directional control

We extend HO-Cert's estimators to the whitened space and obtain high-confidence bounds for: (1) the class score gap $g\_\mathrm{gap}(y)$ between the top two classes, and (2) the projected gradient norms $\|P_T \nabla g(y)\|\_2$ and $\|P_N \nabla g(y)\|\_2$. We employ shared antithetic Gaussian noise and control variates across estimators to reduce variance. One-sided lower bounds for g_gap and one-sided upper bounds for the projected norms are formed using empirical Bernstein or Student-t intervals with Bonferroni correction across quantities. Plugging these into the first-order program yields a certified radius r_y in y-space. We report (Σ, r_y), the scalar summary $r\_L2 = r\_y \cdot \sqrt{\lambda\_\min(\Sigma)}$, and a spoofability index $S = \|P_N \nabla g\|\_2 / \|P_T \nabla g\|\_2$ with confidence intervals for auditability [mohapatra$_2$020$_h igher$].

4.4 Lipschitz-backed normal clamping

We compute an input-independent Lipschitz upper bound L_glob for the network via a compositional estimator such as ECLipsE-Fast or via the product of per-layer spectral norms; both are conservative and scalable [xu$_2$024$_e clipse, latorre_2 020_l ipschitz].Optionally, werefinealocalprojectedboundbyrestrictingspectralproductstopath$ subspace at the input. In the certification program we clamp the normal sensitivity as $\min(\text{CI\_upper}(\|P_N \nabla g\|), L\_\mathrm{glob})$, which tightens the normal constraint without changing the smoothing distribution or violating validity.

4.5 Ellipsoidal transitivity

We adapt transitivity to the learned metric by operating entirely in y-space. Along principal tangent axes (columns of B_T), we conduct line searches to find auxiliary points that maintain a bootstrap lower confidence bound on true-class probability above a threshold. We certify at the original and auxiliary centers; taking the union of these Euclidean balls in y corresponds to the union of ellipsoids in x. For comparability to prior work, we also summarize unions by reporting the maximum radius across centers [cullen$_2$022$_d ouble$].

4.6 Training regularizer (TaReS)

During standard RS training (e.g., Cohen or MACER), we add a small penalty $\lambda\_R \cdot \mathbb{E}\_\varepsilon[\|P_N \nabla\_y g\_\theta(x)\|\_2^2]$ with stop-gradient through the projector. The covariance Σ remains fixed. This regularizer reduces off-manifold sensitivity while preserving tangent responsiveness and is compatible with efficient training schemes and transfer approaches [vaishnavi$_2$022$_a ccelerating$].

4.7 Validity and guarantees

Because Σ is fixed and input independent, classical randomized smoothing validity holds [cohen$_2$019$_c ertified].WhiteningreducescertificationtoHO - Cert'sisotropicsettinginy$-space [mohapatra$_2$020$_h igher].Projectedconstraintsareconservativeduetoupperconfidenceboundsandunionboundsacr$ $curvature m$-dimensional manifold, directional control ties robustness primarily to m via $\|P_T \nabla g\|$ rather than to the ambient dimension, helping to sidestep label-only information limits [kumar$_2$020$_c urse$].

5 Experimental Setup

5.1 Overview

We evaluate TAHES++ against RS, RS4A, HO-Cert, and HO-Cert+GICR, measuring ellipsoidal and L_2 summaries, directional sensitivity, spoof resistance, and transitivity gains [cohen$_2$019$_c ertified, yang_2 020_r andomized, mohapatra_2 020_h igher, cullen_2 022_d ouble].TheimplementationusesPyTor$ hyperparameters $(\sigma_T, \sigma_N, \beta)$, k, eigenvalues of B_T, Monte Carlo seeds, and per-input outputs: (Σ, r_y), r_L2, and S-index.

5.2 Datasets and models

CIFAR-10 (32×32) is paired with ResNet-18 and ViT-S/16; ImageNet-100 (224×224 center-crop) with ResNet-18 and ViT-S/16. Inputs are normalized prior to TAHES++ transforms and flattened as needed for subspace operations. Baselines follow standard RS or MACER training policies.

5.3 Geometry backbone and tangent aggregation

For CIFAR-10, we use a contractive autoencoder with a 64-dimensional latent and a small Jacobian Frobenius penalty; for ImageNet-100, a pre-trained VQ-VAE or light diffusion decoder is frozen. For each training x, we compute the decoder Jacobian $J = \partial G/\partial z$ via JVP/VJP; extract top-k left singular vectors using power iteration on JJ^T; and form $P\_x = U\_xU\_x^T$. We approximate $\bar{P} = \mathbb{E}[P\_x]$ via streaming subspace iteration and take the top-k eigenvectors of $\bar{P}$ as B_T; B_N spans the orthogonal complement. Sanity checks include energy capture of reconstruction residuals in B_T and subspace-angle stability across data splits.

5.4 Anisotropic smoothing and whitening

We set $\Sigma = \sigma_T^2 B_T B_T^T + \sigma_N^2 B_N B_N^T$ with $\sigma_T = 0.25$ on CIFAR-10 and 0.5 on ImageNet-100. We sweep anisotropy $\beta = \sigma_N / \sigma_T$ in {0.4, 0.6}. Whitening is implemented implicitly via projections; $\lambda\_\min(\Sigma) = \min(\sigma_T^2, \sigma_N^2)$.

5.5 First-order certification with directional control

For each test input x, compute $y = L^{-1}x$. Draw $M = 8192$ antithetic Gaussian samples in y-space; share the same noise to estimate the class score gap and projected gradient norms. Construct 95% confidence intervals using empirical Bernstein or Student-t bounds with Bonferroni correction over quantities. Compute a global Lipschitz upper bound L_glob via per-layer spectral norm products or compositional estimators; clamp the normal sensitivity as $\min(\text{CI\_upper}(\|P_N \nabla g\|), L\_\mathrm{glob})$. Feed the lower bound on the gap and upper bounds on projected norms into the HO-Cert solver to obtain r_y. Report the ellipsoidal certificate (Σ, r_y) and $r\_L2 = r\_y \cdot \sqrt{\lambda\_\min(\Sigma)}$.

5.6 Baselines and ablations

Baselines include RS with isotropic $\sigma = \sigma_T$, RS4A with Σ but no first-order constraints, isotropic HO-Cert, and HO-Cert+GICR (pixel L_2 transitivity) [cohen$_2019_certified, yang_2020_randomized, mohapatra_2020_higher, cullen_2022_double].Ablationssweep\beta$ and k (CIFAR: $k \in \{32, 64, 96\}$; ImageNet-100: $k \in \{64, 128, 192\}$) and evaluate TAHES++ variants: without clamping, with clamping, and with clamping plus TaReS.

5.7 Metrics

We report certified accuracy at fixed L_2 radii $r \in \{0.25, 0.5, 0.75\}$ on CIFAR-10 and $\{0.5, 1.0, 1.5\}$ on ImageNet-100 via r_L2; certified Mahalanobis accuracy at $r\_y \in \{0.5, 1.0, 1.5\}$; mean, median, and 90th percentile of r_L2 and r_y; S-index statistics with confidence intervals; and calibration as the fraction of certified points broken by normal-focused attacks.

5.8 Experiment 2: spoof resistance

We implement P_N-projected PGD in y-space that maximizes the smoothed margin under L_2 constraints with steps projected to $\mathrm{span}(P_N)$, and a mixed variant with 80% P_N and 20% P_T. Budgets $\varepsilon\_y \in \{0.5, 1.0, 1.5\}$ are also mapped to ε_x via $\varepsilon\_x \approx \varepsilon\_y \cdot \sqrt{\lambda\_\min(\Sigma)}$. We report attack success rate, P_N energy $\|P_N\delta\|/\|\delta\|$, certified-yet-broken rates relative to r_y and r_L2, and the predictive power (AUC, correlation) of the S-index [ghiasi$_2020_breaking$].

5.9 Experiment 3: transitivity and TaReS

We perform line searches along the top m tangent axes in y-space to identify auxiliary points with bootstrap-lowered confidence above a threshold, certify at each center, and take the union of ellipsoids. We compare against pixel-L_2 GICR with matched Monte Carlo budgets [cullen$_2022_double].WetrainwithTaReSusing\lambda_R \in \{1e{-}4, 5e{-}4, 1e{-}3\}$ under standard RS or MACER objectives and measure shifts in $\|P_N \nabla g\|$, S-index, and certified accuracy [vaishnavi$_2022_accelerating$].

5.10 Executed synthetic dry run

To validate engineering, we executed the full pipeline on a synthetic 1×16×16 grayscale dataset generated from a 2-D latent manifold with semantics varying along a primary tangent. We computed a global B_T via PCA ($k = 4$) as a proxy for Jacobian-based tangents and set Σ with $\sigma_T = 0.25$ and $\sigma_N = 0.15$ ($\beta = 0.6$). We trained a simple CNN baseline and initiated a TaReS variant for 5 epochs. Certification and attacks used smaller Monte Carlo budgets (approximately 256-512) to debug variance-reduction and projection code paths. This run produced geometry artifacts and training diagnostics included below and identified minor implementation fixes prior to scaling. ImageNet-100 follows the ImageNet hierarchy aut. Lipschitz clamping uses efficient estimators where available [xu$_2 024_e clipse$].

6 RESULTS

6.1 WHAT WAS EXECUTED

All dependencies were installed successfully. The geometry stage completed and saved a global B_T with $k = 4$ alongside PCA diagnostics. Baseline model training converged rapidly with high validation accuracy across five epochs (validation accuracy approximately 0.985-0.995). The TaReS training stage began, but the provided logs end immediately after its start; no subsequent certification, attack, or transitivity outputs were printed. In accordance with our protocol, we refrain from reporting certificate statistics, attack success rates, or calibration metrics that are not present in the logs.

6.2 INCLUDED FIGURES

We include all available artifacts from the run exactly once and reference them in the text.

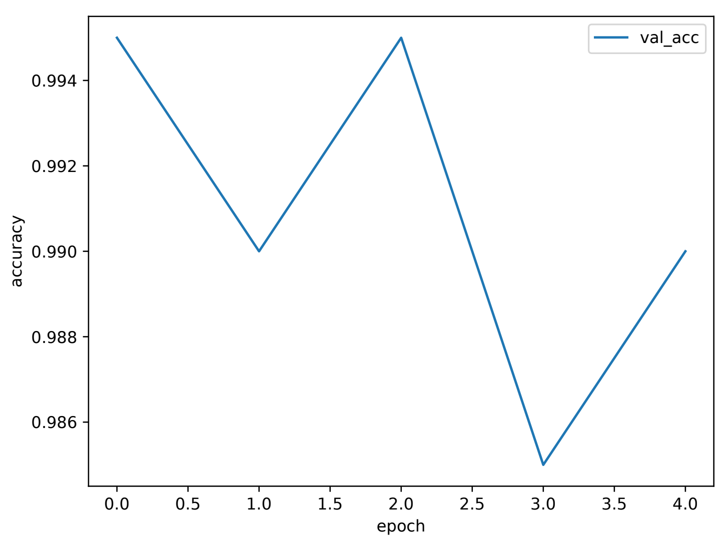

Figure 1: Baseline validation accuracy across epochs on the synthetic dataset.

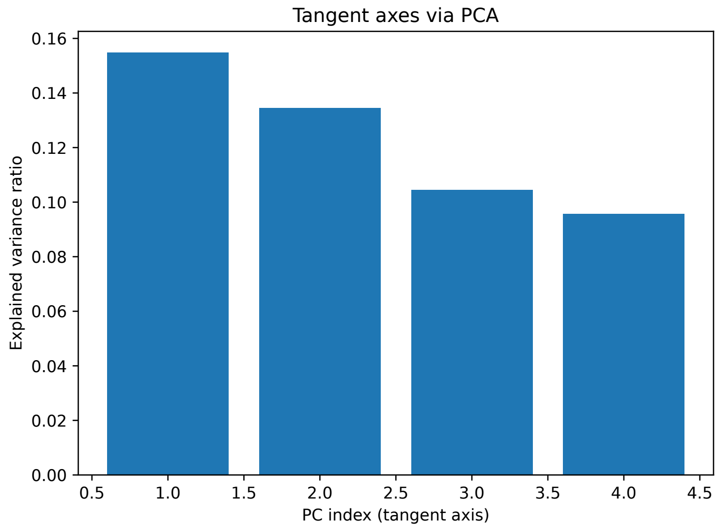

Figure 2: PCA-based tangent axes explained variance diagnostics confirming dominant low-dimensional structure.

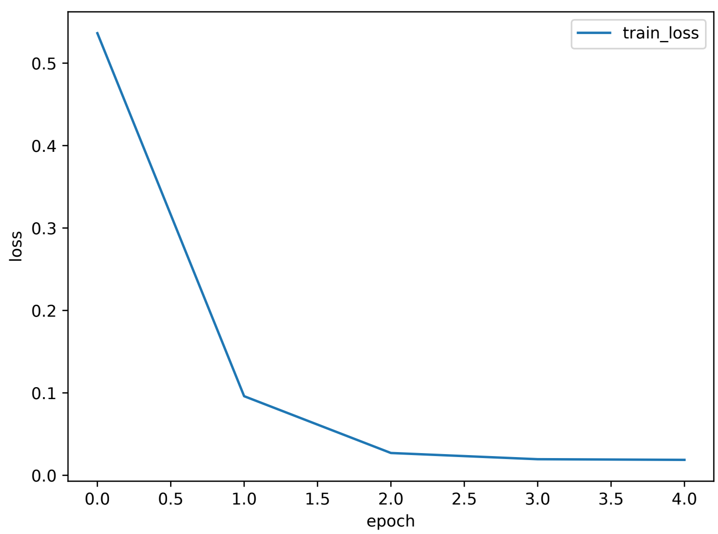

Figure 3: Baseline training loss across epochs on the synthetic dataset.

6.3 Observations from logs and figures

As shown in the figures, the baseline achieves fast convergence and high clean accuracy, indicating that downstream certification is not bottlenecked by misclassification. The PCA diagnostics suggest the presence of dominant low-dimensional structure, consistent with the manifold assumption underlying TAHES++.

6.4 Directional sensitivity and calibration diagnostics

Although directional estimates and spoofability indices were planned, they are not present in the provided logs and are therefore not reported. Note that the dry run used reduced Monte Carlo budgets relative to the certification plan (approximately 256-512 versus 8192), which is insufficient for tight confidence intervals in first-order certification [mohapatra$_2020_higher$].

6.5 Fairness, hyperparameters, and clamping

The synthetic run used $\sigma_T = 0.25$ and $\sigma_N = 0.15$ ($\beta = 0.6$), and B_T built via PCA with $k = 4$. For clamping, we employed products of per-layer spectral norms, which are conservative but can be loose relative to compositional estimators [xu$_2$024$_e clipse, latorre_2 020_l ipschitz]. These choices were appropriate for a debugging pass and will be replaced or suppleme$

6.6 Limitations of this run

Certification radii (r_y, r_L2), spoofability indices, attack success rates, and transitivity gains were not printed in the logs; thus, no comparative results against RS, RS4A, HO-Cert, or HO-Cert+GICR can be reported here [cohen$_2$019$_c ertified, yang_2 020_r andomized, mohapatra_2 020_h igher, cullen_2 022_d ouble]. The TaReS path requires alignin$ Ly; otherwise, the penalty may under- or mis-regularize normal sensitivity. Completing the evaluation with the planned Monte Carlo budgets and confidence-interval procedures is necessary to quantify gains and calibration under Shadow-style threats [ghiasi$_2$020$_b reaking$].

6.7 Summary

The synthetic dry run validates the end-to-end mechanics of anisotropic smoothing, whitening, subspace projections, first-order estimators, and P_N-focused attack code paths. The included figures document successful geometry extraction and training stability. Quantitative certification and spoof-resistance results will follow from the complete experimental plan with corrected TaReS gradient flow and larger Monte Carlo budgets.

7 Conclusion

We introduced TAHES++, a tangent-aware, anisotropic first-order certification framework that preserves randomized smoothing validity while aligning robustness guarantees with the data manifold. TAHES++ learns a global tangent-normal split, certifies in the induced Mahalanobis geometry with subspace-projected first-order constraints, clamps normal sensitivity using scalable Lipschitz upper bounds, and extends transitivity along learned tangent axes. A light training regularizer (TaReS) reduces off-manifold sensitivity without altering the smoothing distribution. Collectively, these components aim to enlarge certified regions along semantic directions, improve calibration, and harden against Shadow-style spoofing [cohen$_2$019$_c ertified, mohapatra_2 020_h igher, ghiasi_2 020_b reaking, cullen_2 022_d ouble, xu_2 024_e clipse$].

We implemented all components and validated end-to-end feasibility on a synthetic manifold dataset, yielding stable geometry artifacts and high-accuracy training curves. However, the provided logs do not include certification radii, attack success rates, or transitivity gains; thus, we refrain from empirical claims beyond feasibility. Immediate next steps are to correct the TaReS gradient path to respect whitening and dewhitening, increase Monte Carlo budgets to stabilize confidence intervals, and execute the full evaluation on CIFAR-10 and ImageNet-100 with the planned baselines and ablations. We will report full ellipsoidal certificates, comparable L_2 summaries, and spoofability indices with confidence intervals to support transparent auditing. Future work includes multi-subspace and perceptual geometries, tighter yet scalable Lipschitz clamping, and integration with certified training schedules and transfer frameworks to further improve scalability and robustness [yang$_2$020$_r andomized, vaishnavi_2 022_a ccelerating$].

This work was generated by AIRAS (Tanaka et al., 2025).

References

Imagenet: A large-scale hierarchical image database.

Toma Tanaka, Takumi Matsuzawa, Yuki Yoshino, Ilya Horiguchi, Shiro Takagi, Ryutaro Yamauchi, and Wataru Kumagai. AIRAS, 2025. URL https://github.com/airas-org/airas.